Spatial distance and curvature from an exact counting identity in two-dimensional causal sets

Roberto Bada[1,2]

[1] SapereIA, Barcelona, Spain

[2] Institute of Cosmos Sciences, University of Barcelona (ICCUB), Barcelona, Spain

E-mail: roberto@sapereia.com

Abstract. In causal set theory, geometry must be recovered from the causal order and from counting. In two dimensions, the causal-overlap regions of Boguñá and Krioukov satisfy an exact identity: for two spacelike elements and *any* element c in their common past, $d^2 = 2V_AV_B/V_C$ in flat spacetime. With a Poisson bias correction, the identity gives a spatial-distance estimator that is unbiased for d^2 above two discreteness lengths. On the same pairs, this estimator is more accurate than the equal-depth construction of Boguñá and Krioukov at all separations, and the freedom in the choice of c is what improves on it. Between two and ten discreteness lengths it is as accurate as, or more accurate than, the past–future construction of Hu, Meyer and Meyer, which is more accurate at larger separations. In curved spacetime the identity fails in a way that depends on the depth of c. In de Sitter space $d^2_{\rm est}/d^2 = \tanh^2(H\tau/2)/[2\ln\cosh(H\tau/2)] = 1 - R\tau^2/16 + O(\tau^4)$, and in any $1+1$ FRW spacetime, to leading order in the depth, the failure measures the scalar curvature one third of the way from the pair to c. A fit of this depth drift, with its coefficients summed over the pairs of many independent sprinklings, reproduces the same fit on the exact continuum volumes. It does so for pairs inserted by hand (de Sitter, flat and FRW sprinklings) and for all sprinkled pairs of a position-based selection (de Sitter and flat, within 1.7σ); selecting the elements nearest to given points biases it low. On a sprinkling of a causally convex region, such as an interval of the causal set, the fit window needs no coordinates. For the curvature of a given region, however, the fit is less precise than the ordering fraction (in de Sitter space its spread between sprinklings is about three times larger) and, within one region, biased low. As a supporting diagnostic, we calibrate the Lorentzian-embedding test of Clough and Evans and add the fraction of null-separated pairs, which lets the test reject layered orders on its own. An appendix gives a Laguerre identity for the mean of Johnston's propagator at fixed interval size.

Keywords: causal sets, two-dimensional orders, spatial distance, scalar curvature, de Sitter space, Lorentzian embedding

Submitted to: *Class. Quantum Grav.*

1. Introduction

A causal set is a locally finite partially ordered set: every interval $I(x, y) = \{z : x \prec z \prec y\}$ is finite [1, 2]. Causal set theory proposes that spacetime is such a set, with the continuum emerging at large scales. The number of elements in a region then measures its spacetime volume, and the order fixes the light cones. The kinematic program is to recover all other geometric information from these two ingredients.

Part of this program is well established. The longest chain between two elements estimates their proper time [3]; interval and chain abundances give the dimension [4, 5] and the curvature [6, 7]; and the Benincasa–Dowker operator gives the scalar curvature and the action [8, 9]. Spatial distance is harder, because spacelike elements bound no interval. Boguñá and Krioukov measured it through the overlap of the pasts of two elements, as seen from an element c at equal proper time below both [10]. Brightwell and Gregory, and Rideout and Wallden, used the smallest interval containing both elements [3, 11]. Hu, Meyer and Meyer gave an exact reconstruction in $1 + 1$ dimensions from a past and a future element [19]. Whether a causal set resembles a manifold at all is tested by dimension estimators, interval abundances [7] and Lorentzian embeddings [20, 22, 23]. Matter can be introduced through Johnston's sum over chains [24, 25], whose average over sprinklings reproduces the Klein–Gordon propagator.

Two-dimensional causal sets are a natural laboratory. Every $1 + 1$ metric is conformally flat, intervals are rectangles in light-cone coordinates, and a sprinkled diamond is a uniformly random two-dimensional order, whose combinatorics is known [26, 27]; the $\mathrm{e}^{-\beta S}$ model of two-dimensional orders also has a well-studied phase structure [26, 27, 28]. This paper is built on one exact identity that follows from these facts, and on the use of its violation as a measure of curvature. The *rectangle identity* expresses the spatial distance of two elements through the overlap volumes seen from any element of their common past (section 3). It removes the equal-depth requirement of [10] and yields a distance estimator, which we compare with the equal-depth, past–future and Hu–Meyer–Meyer constructions on the same pairs. In curved spacetime the identity fails as a function of the depth of c. We derive this depth drift in closed form in de Sitter space and to leading order in any $1 + 1$ FRW spacetime, and turn it into a curvature estimator (section 4). Spatial distance in causal sets has also been approached through thickened antichains and induced distances on spatial hypersurfaces [12, 13, 14]. Curvature has also been obtained from interval and chain abundances in small causal diamonds [6, 15] and from the Benincasa–Dowker operator [8]; we compare our curvature estimator with the last two at equal density.

As a supporting diagnostic, we calibrate the Lorentzian-embedding test of Clough and Evans and add a statistic that lets it reject layered orders on its own (section 5). We apply it to sequential growth and to configurations of the $\mathrm{e}^{-\beta S}$ model (section 6). Section 7 discusses novelty and limitations. Appendix A gives a combinatorial identity for Johnston's propagator at fixed interval size, and Appendix B summarizes other approaches we tested. Spatial distances in $3+1$ dimensions, where the rectangle identity

has no analogue, are treated in a companion paper [29].

2. Setting

Test causal sets are generated by sprinkling: a Poisson process of density ρ in a region of a given spacetime, with the order induced by the light cones. The coordinates are then discarded and used only for comparison. We write $\ell = \rho^{-1/2}$ for the discreteness length: a region of volume ℓ^2 contains one element on average.

We use light-cone coordinates $u = t - x$, $v = t + x$, in which $\mathrm{d}t\,\mathrm{d}x = \mathrm{d}u\,\mathrm{d}v/2$, and $x \prec y$ if and only if $u_x < u_y$ and $v_x < v_y$. The interval between $x \prec y$ is the rectangle $[u_x, u_y] \times [v_x, v_y]$, with volume $V = \Delta u\,\Delta v/2 = \tau^2/2$. A sprinkled causal diamond is therefore a uniformly random two-dimensional order, the intersection of two random linear orders [26, 27]. For two spacelike elements, $d^2 = \Delta x^2 - \Delta t^2 = -\Delta u\,\Delta v$.

The proper time obtained from the longest chain has mean $\simeq 2\sqrt{n}$ in an interval of n elements, with Tracy–Widom fluctuations of order $n^{1/6}$ [30]. Wherever a volume is available, we use the count instead: $\tau^2 = 2n/\rho$. Pairs of elements in one sprinkling share elements, so their errors are correlated; quoted uncertainties are therefore computed from the spread between independent sprinklings.

3. The rectangle identity and a distance estimator

Let a, b be spacelike and let c be an element in their common past. Following [10], we split the two intervals into

$$A = I(c,a) \setminus I(c,b), \qquad B = I(c,b) \setminus I(c,a), \qquad C = I(c,a) \cap I(c,b). \tag{1}$$

Boguñá and Krioukov choose c at equal proper time below a and b, and derive d from the overlap fraction. In $1+1$ dimensions the choice of c is immaterial, because of the following identity (figure 1).

Derivation. Place c at the origin and take $u_a < u_b$ and $v_a > v_b$ (a and b spacelike). Then $I(c,a) = [0,u_a] \times [0,v_a]$ and $I(c,b) = [0,u_b] \times [0,v_b]$. The rectangle from c to the join (u_b, v_a) splits into four tiles,

$$\begin{aligned} C &= [0,u_a] \times [0,v_b], \quad & A &= [0,u_a] \times [v_b,v_a], \\ B &= [u_a,u_b] \times [0,v_b], \quad & D &= [u_a,u_b] \times [v_b,v_a]. \end{aligned} \tag{2}$$

With $\Delta u = u_b - u_a$ and $\Delta v = v_a - v_b$, the volumes are $V_C = u_a v_b/2$, $V_A = u_a \Delta v/2$, $V_B = \Delta u\, v_b/2$ and $V_D = \Delta u\,\Delta v/2$, so that $V_C V_D = V_A V_B$. The tile D has a and b at opposite corners, so $V_D = d^2/2$. Hence, for every c in the common past,

$$d^2 = \frac{2 V_A V_B}{V_C}. \tag{3}$$

The tile D is the interval between the meet and the join of a and b, i.e. the region whose volume defines the $1+1$ spatial distance of Brightwell and Gregory and of Rideout and

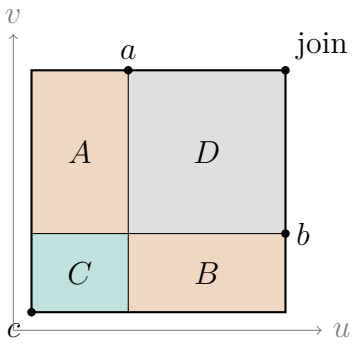


Figure 1. The rectangle identity in light-cone coordinates (rotated by 45° with respect to the usual picture). The rectangle from c to the join of a and b splits into the overlap C, the strips A and B, each seen by only one of a and b, and the tile D, which has a and b at opposite corners. Opposite tiles have equal products, $V_C V_D = V_A V_B$, and $V_D = d^2/2$.

Wallden [3, 11]. The identity shows that this volume can be read from the past alone, through a single element c. The equal-time choice of [10] corresponds to $V_A = V_B$, in which case their $1 + 1$ formula reduces to equation (3). For general c, the search for equal-time elements, which requires the density and the dimension, is unnecessary. The identity relies on the overlap of two $1 + 1$ intervals being again an interval. In higher dimensions it is not, and equation (3) has no direct analogue.

Estimator. On a causal set, V_X is replaced by N_X/ρ, where N_X is the number of elements in the region X. The three regions are disjoint, so N_A, N_B and N_C are independent Poisson variables with means $\lambda_X = \rho V_X$. The product $N_A N_B$ is unbiased for $\lambda_A \lambda_B$, but $1/N_C$ is biased and is undefined at $N_C = 0$. Since $E[1/(N + 1)] = (1 - \mathrm{e}^{-\lambda})/\lambda$ for a Poisson variable N with mean λ,

$$\widehat{d^2} = \frac{2 N_A N_B}{\rho\,(N_C + 1)} \tag{4}$$

is unbiased for d^2 up to a factor $1 - \mathrm{e}^{-\lambda_C}$. The relative variance is approximately $1/\lambda_A + 1/\lambda_B + 1/\lambda_C$, so the deepest available c is optimal, and we take the element with the largest N_C. The strips A and B are long and thin and contain many elements even when $d \sim \ell$; this is why the estimator works down to the discreteness scale. By contrast, a direct count of the tile D has relative variance $1/\lambda_D = 2\ell^2/d^2$, which is $\frac{1}{2}$ at $d = 2\ell$.

Tests. We use two pair selections, which differ in the size of the errors. Table 1 and figure 2 use random spacelike pairs among *late* elements of flat sprinklings of a unit diamond; these pairs have large common pasts. Table 2 uses pairs in the *middle* of the diamond, which have both a common past and a common future, but smaller ones. In both cases, c is the deepest element for each pair. The median error falls approximately as $N^{-1/2}$ at large d. Below 2ℓ, the discreteness scale itself, it remains at 15–25%, and $\widehat{d^2}$ is also biased high there [by $(37 \pm 7)\%$ at $N = 4000$ on the pairs of table 2].

Table 2 compares the estimator with the existing constructions on the same pairs: the equal-depth filter of Boguñá and Krioukov [equation (34) of [10]]; the smallest interval $I(c, f)$ with $c \prec a, b \prec f$ (Brightwell–Gregory, Rideout–Wallden), $d^2 = 2(N_{\min} - 2)/\rho$; and the construction of Hu, Meyer and Meyer [19]. They showed that for $p \prec a, b \prec q$ with all four proper times equal to s, $d^2 = \tau_{pq}^2 - 4s^2$, and gave this as an exact statement about Minkowski space. On the causal set, we read every τ^2 from a count ($\tau^2 = 2N/\rho$), take p from the common past and q from the common future, and keep the pairs (p, q) whose four counts agree within Poisson noise ($\chi^2 < 3$). Among these we use the pair with the smallest counts, which has the smallest variance; choosing instead the best-matched pair, or averaging over all matched pairs, is less accurate. Two ablations separate the ingredients of our estimator: the same deepest c without the Poisson correction, $2N_A N_B/(\rho N_C)$, and the equal-depth elements c of [10] with the corrected formula, equation (4).

Above 2ℓ the rectangle estimator is unbiased for d^2 within the uncertainties. Between 2ℓ and 10ℓ its median error is the smallest. The Hu–Meyer–Meyer estimate is close in error but biased high by about 20%, because $\tau_{pq}^2 - 4s^2$ is a small difference of two noisy terms. Above 10ℓ the Hu–Meyer–Meyer estimate is unbiased and the most accurate, with half the error of ours or less. There the tile between a and b is large, and a direct count of it, corrected exactly by the four equal intervals, beats reading the same volume from the strips. For d itself the rectangle estimator reads 1–3.5% low, as expected from the square root of an unbiased estimate of d^2. The plain past–future count is biased high, because the smallest interval between elements cannot coincide with the tile D.

The ablations locate the gain over the equal-depth estimator. With the deepest c, N_C is large and the Poisson correction makes no difference. With the equal-depth elements, where N_C is small, the correction removes the bias of the equal-depth estimator ($+10\% \to +1\%$ and $+6\% \to +2\%$ between 2ℓ and 10ℓ) but not its larger error (0.129 and 0.104, against 0.095 and 0.059 with the deepest c). The bias of [10] on the causal set thus comes from the uncorrected ratio of small counts, and the gain in accuracy comes from the freedom to use the deepest c, which the identity provides. The equal-depth filter also returns no estimate for 10–30% of the pairs. Below 2ℓ every estimator is poor, and the two past–future constructions fail entirely (median errors above 50%). For the baseline tests we add a bottom element that precedes all others. The deepest c is then always that element; it is selected without reference to its own count, so the selection introduces no bias. (In table 1, which has no such element, the deepest c is one of the lowest elements.) The freedom in the choice of c is exercised by the curvature tests of section 4.

4. Curvature from depth drift

In flat spacetime every c in the common past gives the same estimate. In curved spacetime equation (3) no longer holds exactly, and the estimate depends on the depth

Table 1. Median relative error of equation (4) with the deepest c; flat 1+1 sprinklings, 300 pairs among late elements in each of 4 sprinklings per N. (Table 2 uses pairs in the middle of the diamond, with smaller common pasts, and therefore has larger errors.)

N	$d < 2\ell$	2ℓ–10ℓ	$> 10\ell$
1000	23%	6.5%	5.2%
2000	15%	5.0%	3.4%
4000	19%	5.1%	2.5%

Table 2. Rectangle estimator, equation (4), against the equal-depth filter of [10], the past–future construction of Hu, Meyer and Meyer [19] read from counts (see text), and the smallest past–future interval [3, 11], on the same spacelike pairs in the middle of flat diamonds (12 sprinklings × 100 pairs per N). Two ablations: the deepest c without the Poisson correction, and the equal-depth elements c with it. Bias: mean relative error of d^2, in percent, with the standard error from the spread between sprinklings in parentheses. Error: median absolute relative error of d. The equal-depth rows have no estimate for 10–30% of the pairs.

		2ℓ–10ℓ		$> 10\ell$	
N	method	bias	error	bias	error
1000	rectangle, equation (4)	+1(5)	0.095	−4(4)	0.082
	no correction	+3(5)	0.097	+1(4)	0.080
	equal depth [10]	+10(4)	0.131	+13(3)	0.126
	with correction	+1(3)	0.129	−2(3)	0.127
	Hu–Meyer–Meyer	+23(3)	0.098	+4(2)	0.028
	past–future	+86(3)	0.272	+25(3)	0.095
4000	rectangle, equation (4)	0(2)	0.059	−2(1)	0.039
	no correction	0(2)	0.059	0(2)	0.037
	equal depth [10]	+6(2)	0.106	+2(1)	0.075
	with correction	+2(2)	0.104	−5(1)	0.073
	Hu–Meyer–Meyer	+22(4)	0.099	0(1)	0.015
	past–future	+79(5)	0.282	+14(1)	0.049

of c. This dependence measures curvature.

de Sitter volumes. Two-dimensional de Sitter space in conformal coordinates,

$$\mathrm{d}s^2 = \frac{-\mathrm{d}\eta^2 + \mathrm{d}x^2}{(H\eta)^2}, \qquad \eta < 0, \tag{5}$$

has the light cones of flat space. The tiles A, B, C, D are the same rectangles in $(u, v) = (\eta - x, \eta + x)$, but volumes carry the conformal weight,

$$\mathrm{d}V = \frac{\mathrm{d}\eta\,\mathrm{d}x}{(H\eta)^2} = \frac{2\,\mathrm{d}u\,\mathrm{d}v}{H^2(u+v)^2}. \tag{6}$$

When this is integrated over a rectangle $[u_1, u_2] \times [v_1, v_2]$, the inner integral gives $1/(u + v_1) - 1/(u + v_2)$ and the outer integral gives a logarithm,

$$V_\square = \frac{2}{H^2} \ln \frac{\eta_L\, \eta_R}{\eta_{\mathrm{bot}}\, \eta_{\mathrm{top}}}, \tag{7}$$

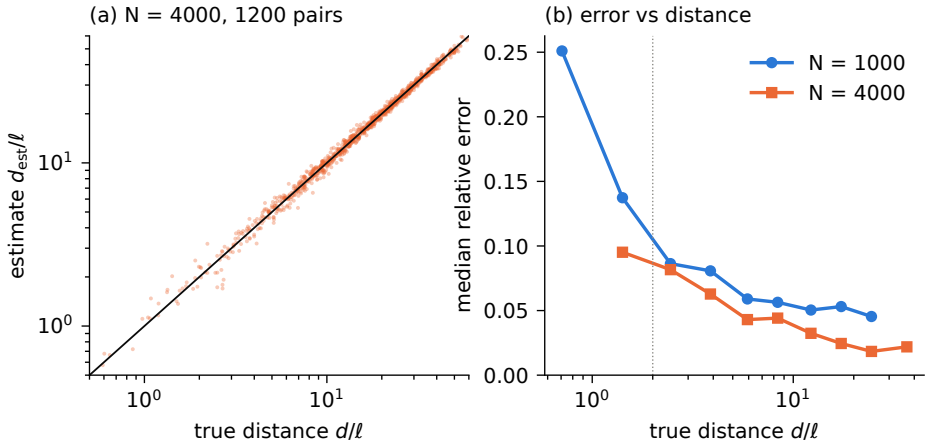


Figure 2. Equation (4) with the deepest c on flat 1+1 sprinklings, computed from the causal order only. (a) Estimate against true distance, $N = 4000$. (b) Median relative error per distance bin; the dotted line marks 2ℓ.

with $\eta = (u + v)/2$ evaluated at the left, right, bottom and top corners (all $\eta < 0$, so the argument is positive).

Closed form. Place a, b on one slice at $(\eta_0, \mp h)$ and c directly below their midpoint at $(\eta_0 - T, 0)$, and set $H = 1$. By symmetry, $V_A = V_B$. Write $\bar{\eta} = |\eta_0|$ and $q = \bar{\eta} + T/2$, so that $|\eta|$ at the corners of A is $q \mp h/2$, $\bar{\eta} + h$ and $\bar{\eta}$. To first order in h, equation (7) gives

$$V_A = 2 \ln \frac{(q - \frac{h}{2})(\bar{\eta} + h)}{(q + \frac{h}{2})\,\bar{\eta}} \simeq 2h\Big(\frac{1}{\bar{\eta}} - \frac{1}{q}\Big) = \frac{hT}{\bar{\eta}\, q}, \tag{8}$$

and at $h = 0$

$$V_C = 2 \ln \frac{q^2}{(\bar{\eta} + T)\,\bar{\eta}} = -2 \ln(1 - x), \qquad x = \frac{T^2}{4q^2}. \tag{9}$$

The true distance is $d \simeq 2h/\bar{\eta}$, so

$$\frac{d_{\text{est}}^2}{d^2} = \frac{2V_A^2/V_C}{4h^2/\bar{\eta}^2} = \frac{x}{-\ln(1 - x)}. \tag{10}$$

The proper time from c to the midpoint of a, b is $\tau = \ln[(\bar{\eta} + T)/\bar{\eta}]$, so that $T/\bar{\eta} = \mathrm{e}^\tau - 1$ and $x = [(\mathrm{e}^\tau - 1)/(\mathrm{e}^\tau + 1)]^2 = \tanh^2(\tau/2)$. Restoring H,

$$\begin{aligned} \frac{d_{\text{est}}^2}{d^2} &= \frac{\tanh^2(H\tau/2)}{2 \ln \cosh(H\tau/2)} = 1 - \frac{(H\tau)^2}{8} + O(\tau^4) \\ &= 1 - \frac{R\,\tau^2}{16} + \ldots, \qquad R = 2H^2, \end{aligned} \tag{11}$$

which is exact as $d \to 0$ at any depth, for c directly below the midpoint of a and b. The denominator is $H^2/2$ times the de Sitter volume of an interval of height τ, $V = (4/H^2) \ln \cosh(H\tau/2)$, and the result is independent of η_0, as required by the de Sitter symmetry. We verified equation (11) against the exact volumes of equation (7)

to 10^{-5} for $H\tau = 0.1$–3. A sideways offset of c changes it little, and a finite d adds an approximately depth-independent shift. We call $1 - d_{\rm est}^2/d^2$ the depth drift.

Other spacetimes. Every $1+1$ metric is conformally flat, so the leading term should hold locally in any $1+1$ spacetime. We checked this in FRW spacetimes $\mathrm{d}s^2 = a(\eta)^2(-\mathrm{d}\eta^2 + \mathrm{d}x^2)$ with $a = \eta^p$, where $R = 2(\ln a)''/a^2 = -2p\,\eta^{-2-2p}$ varies in time and the rectangle volumes have closed forms. For $p = 1$, -2 and $\frac{1}{2}$ (negative, positive and negative curvature), $16(1 - d_{\rm est}^2/d^2)/\tau^2$ tends to R at the location of a, b as $\tau \to 0$ (to 1% at $\tau = 0.01$), with the correct sign: for $R < 0$ the ratio $d_{\rm est}^2/d^2$ exceeds 1. At finite depth the drift measures a curvature between the values at a, b and at c, with more weight near the top.

The weighting follows from the next order. Write $a = \mathrm{e}^{\sigma(\eta)}$, place a, b at $(\eta_0, \mp h)$, and place c a coordinate time T straight below. As $h \to 0$, the strips have $V_A = V_B = h\int_0^T \mathrm{e}^{2\sigma(\eta_0 - s/2)}\mathrm{d}s$, the overlap is $V_C = \frac{1}{2}\int_0^T\int_0^T \mathrm{e}^{2\sigma(\eta_0-(s+r)/2)}\mathrm{d}s\,\mathrm{d}r$, and $d = 2h\mathrm{e}^{\sigma(\eta_0)}$. We expand σ in a Taylor series at η_0, integrate term by term, and trade T for the proper time $\tau = \int_0^T \mathrm{e}^{\sigma(\eta_0 - s)}\mathrm{d}s$. This gives

$$1 - \frac{d_{\rm est}^2}{d^2} = \frac{\tau^2}{16}\Big[R - \frac{\tau}{3}\frac{\mathrm{d}R}{\mathrm{d}\tau} + O(\tau^2)\Big], \tag{12}$$

with $R = 2\sigma''\mathrm{e}^{-2\sigma}$ and its derivative along the worldline (toward the future) evaluated at a, b. To this order, the drift measures R one third of the way from a, b down to c. Equation (12) fixes only this first moment of the depth weighting. A triangular weight $2(\tau - s)/\tau^2$ on the depths $s \in [0, \tau]$ is one illustration with this mean depth, not a kernel determined by the expansion; the next order involves R^2 and $\mathrm{d}^2R/\mathrm{d}\tau^2$ and is not of the form of a weighted average of R. The exact power-law volumes confirm equation (12). For $p = 1$ at $\eta_0 = 1.5$, $16(1 - d_{\rm est}^2/d^2)/\tau^2 = -0.3982$, -0.4073 and -0.4200 at $\tau = 0.02$, 0.05 and 0.1, against -0.3997, -0.4068 and -0.4185 from its two terms ($R = -0.395$ at a, b); the other two values of p agree equally well.

Test on sprinklings. We sprinkle de Sitter space ($H = 1$, $\rho = 10^4$) and, as a control, flat space in the same coordinates. We insert a, b by hand $\approx 14\ell$ apart and evaluate equation (4) for every sprinkled c in their common past such that the smallest rectangle containing $I(c, a)$ and $I(c, b)$ lies in the region. This cut is stricter than the window of the fit below, which requires only the two intervals to lie in the region. Figure 3 shows the mean ratio per depth bin over 400 sprinklings against the continuum value computed for the same elements c. The continuum value includes the effects of finite d and of sideways offsets, and therefore lies slightly below equation (11). The flat control remains at 1, and de Sitter follows the continuum. Neighboring depth bins share elements. Computed from the spread over the 400 sprinklings, the correlation of adjacent bins is 0.80, 0.81, 0.90, 0.94 and 0.96 in de Sitter space (from the shallowest pair of bins to the deepest) and 0.80–0.97 in flat space; bins two apart are still correlated at 0.5–0.9. A run of deviations in adjacent bins therefore counts as a single fluctuation, and the points of figure 3 must not be combined as independent. In the shallowest bin the tile C contains few elements, so the factor $1 - \mathrm{e}^{-\lambda_C}$ of equation (4) is not negligible, and the estimator reads low in both geometries.

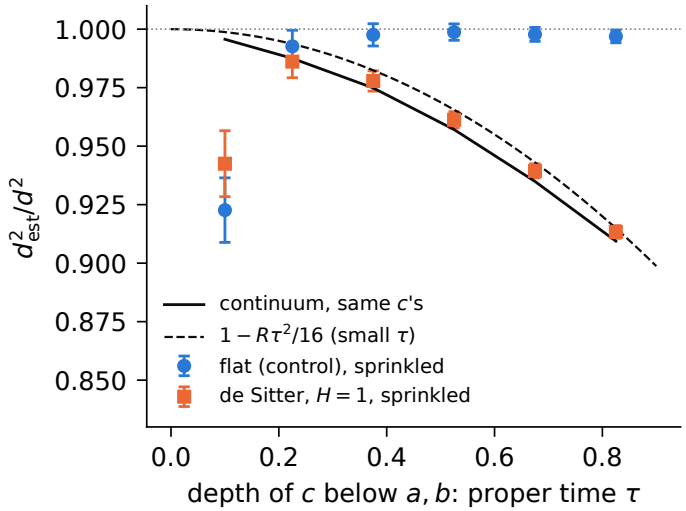


Figure 3. Depth drift of the rectangle estimate: mean over 400 sprinklings ($\rho = 10^4$, $H = 1$) per bin of the proper time τ from c to a, b, with errors from the spread between sprinklings. The bins are strongly correlated (adjacent bins 0.8–0.97; see text), so the error bars are not independent. Solid: continuum ratio from the exact volumes, averaged over the same elements c (finite d and offsets included). Dashed: small-τ term of equation (11).

A curvature estimator. For one pair a, b we fit $\widehat{d^2} = \alpha + \beta\tau^2$ over the elements c of their common past, with $\tau^2 = (N_{I(c,a)} + N_{I(c,b)})/\rho$ from counts, and read off $R = -16\beta/\alpha$. Two choices need care: the fit window and the way pairs are combined.

Window from the order. The fit uses every c with $N_C \geq 20$, which keeps the factor $1 - \mathrm{e}^{-\lambda_C}$ of equation (4) negligible, down to the deepest c whose intervals $I(c, a)$ and $I(c, b)$ lie in the sprinkled region. Neither limit refers to H. The estimator counts only A, B and C, which lie in $I(c, a) \cup I(c, b)$; the tile D is never counted. In a causally convex region that contains c, a and b, these intervals lie in the region. On a sprinkling of a causally convex region, for example an interval of the causal set, every c in the common past can therefore be used, and the window needs no coordinates. Otherwise the boundary must first be inferred, which in $1 + 1$ dimensions is possible from the order [31]. Convexity is an assumption about the sprinkled region: the order decides it only when the causal set is itself defined from the order, as an interval is. Our main tests use a coordinate box, where we decide the condition from the coordinates. We therefore repeated them in a causally convex region, the past of the point $(-0.95, 0)$ cut at $\eta = -3$ (convex, although the cut is not defined by the order), using every c of the common past (1000 sprinklings each). The results agree with table 3. In de Sitter space the pooled estimate (defined below) is 2.10 ± 0.09 and 2.03 ± 0.06 with inserted pairs, and 2.10 ± 0.08 and 2.04 ± 0.06 with all pairs in the windows (see below), at $\rho = 4000$ and 10^4 respectively, against continuum values of 2.03 and 2.04 (the deeper region changes the target; see below). In flat space it is 0.01 ± 0.02 and 0.00 ± 0.02 at $\rho = 4000$, and 0.00 ± 0.01 and -0.01 ± 0.01 at $\rho = 10^4$.

Because this window reaches $H\tau \approx 1$, the τ^4 term of equation (11) contributes. The same fit applied to the exact continuum volumes of the same elements gives $R_{\rm fit}^{\rm cont} = 2.07$–$2.08$ in the de Sitter box (and 2.03–2.04 in the convex region) instead of 2. The estimator reproduces this value, not 2. Since the value is not known in practice, the difference is a systematic error of the method, of +2 to +4% here, which depends on the size of the region. We varied the lower cut ($N_C \geq 5$, 20, 50) and the upper cut (τ^2 up to 50%, 75% or 100% of its largest value). These changes move the pooled de Sitter result only between 1.91 and 2.01 at $\rho = 10^4$ (flat: -0.01 to 0.00), within 0.06–0.28 of the continuum value of each rule; the largest offset belongs to the shortest window, which is also the noisiest.

Pooling. A single pair gives a noisy estimate (standard deviation 3.5 at $\rho = 4000$ and 2.4 at $\rho = 10^4$ in de Sitter space), and the mean of single-pair values is biased low. The reason is that the intercept and slope of one least-squares fit have anticorrelated errors, so $E[-16\beta/\alpha]$ is smaller than $-16E[\beta]/E[\alpha]$. Summing the fit coefficients over the pairs before forming the ratio, $R_{\rm est} = -16\sum\beta/\sum\alpha$, removes this ratio bias to first order. Table 3 shows the pooled estimator over 1000 sprinklings, with errors from a bootstrap over sprinklings. With inserted pairs it agrees with the continuum value of the same fit: in flat space within the errors, and in de Sitter space within 0.9σ and 1.1σ (offsets of -0.095 ± 0.111 and -0.081 ± 0.072 before rounding).

Sprinkled pairs. How a and b are chosen among the sprinkled elements matters. Taking the elements nearest to two target points selects elements with an empty neighborhood (the nearest point lies on the boundary of an empty disc), and part of that disc lies at the shallow ends of the strips A and B. The counts confirm this: averaged over the elements c of the window, $N_A - \rho V_A$ is lower by 0.2–0.4 elements than for inserted pairs, at every depth. Since λ_A grows with the depth of c, a fixed deficit lowers $\widehat{d^2}$ most for shallow c. This tilts the fitted slope upward and gives a low R, in flat space as well; the effect shrinks with density. With this selection the estimator reads low in de Sitter space by 0.37 ± 0.11 at $\rho = 4000$ and by 0.24 ± 0.07 at $\rho = 10^4$ (table 3).

The unbiased use of sprinkled pairs follows from the Mecke formula [33]: a sum over all pairs of points of a Poisson process, with a weight that depends only on the positions of the pair and not on the rest of the configuration, equals the average over pairs placed by hand at the same positions. The requirement is therefore to sum over all pairs of a selection that does not depend on the counts used by the estimator, optionally thinned at random with a fixed probability, and to form the ratio from the pooled sums. We therefore take every spacelike pair with a in a square window of half-width 0.025 ($\approx 2.5\ell$ at $\rho = 10^4$) around one target point and b in the window around the other. We keep each pair with a fixed probability, chosen so that about 20 pairs per sprinkling remain, and pool them all. This removes most of the shortfall. In de Sitter space the pooled estimate is 1.91 ± 0.10 and 2.02 ± 0.06 against continuum values of 2.08 (1.7σ and 1.1σ below, against 0.9σ and 1.1σ for inserted pairs), and in flat space it is 0.00 ± 0.02 and -0.01 ± 0.01 (table 3); in the causally convex region it is 2.10 ± 0.08 and 2.04 ± 0.06 against 2.03 and 2.04. Drawing a single pair uniformly from the two windows gives

results in between (1.80 ± 0.11 and 1.90 ± 0.07): choosing one element out of the N_W in a window weights each configuration by $1/N_W$, which removes about one element from the window. In practice, therefore, the estimator must be pooled over all pairs of a position-based selection, not over a single representative chosen per region. A selection by count-defined separation does not satisfy the condition, because it depends on the same counts as the estimator; we have not tested it.

FRW spacetime. We repeated the test in the $p = 1$ FRW spacetime ($\rho = 10^4$, 1000 sprinklings), with a, b inserted 14ℓ apart at $\eta_0 = 1.5$, using the pooled estimator and the window of the previous paragraphs. The region is causally convex (the past of an element at $\eta = 1.55$) and is cut at three depths. The resulting windows reach $\tau = 0.52$, 0.70 and 0.84, over which R ranges from -0.40 at a, b to -1.37, -2.79 and -6.32. The pooled estimate is -1.03 ± 0.37, -0.88 ± 0.18 and -0.93 ± 0.11, against -0.70, -0.79 and -0.90 for the same fit on the exact volumes of the same elements: the estimator reproduces its continuum target within 1σ. That target is not the leading-order prediction of equation (12), which gives -0.52, -0.55 and -0.59 when fitted in the same way. For the middle window the difference splits about evenly between the higher orders in τ (the exact volumes with c straight below and $d \to 0$ give -0.67) and the finite separation of a and b (with $d = 14\ell$ the same curve gives -0.80). Over a realistic window, equation (12) therefore gives the direction of the effect, a curvature measured below the pair, but falls short of its size by 26–35% in the three windows. The size is set by the exact continuum fit, which the counts reproduce.

Comparison with other estimators. The Benincasa–Dowker operator [8], in its nonlocal two-dimensional form with smearing parameter ϵ, applied to a constant field at the midpoint of the pair, gives $R = -2B_\epsilon[1]$. On the de Sitter and flat sprinklings above, its single-point standard deviation at $\rho = 4000$ is $\approx 1.2 \times 10^3$, 160 and 23–30 for $\epsilon = 0.1$, 0.03 and 0.01, respectively, and it is larger at $\rho = 10^4$ (72–82 at $\epsilon = 0.01$), as expected at fixed ϵ. Stronger smearing lowers the spread (2.7–4.5 at $\epsilon = 0.003$), but the operator then needs the past light cone far beyond our sprinkled region, and the truncation dominates its mean (+0.7 in flat space, +4.7 in de Sitter space). A fair comparison at small ϵ needs a much larger region, which we have not run. The purpose of this operator is the action rather than a pointwise curvature estimate.

The relevant competitor is the ordering fraction of a causal diamond, the leading chain-abundance route [6]. We compare the two as they would be used: one sprinkled region, one estimate. The region is a causal diamond (an interval of the causal set) in de Sitter space, $u, v \in [-3, -0.95]$ ($H = 1$; 2.5×10^3 elements at $\rho = 4000$ and 6.3×10^3 at $\rho = 10^4$), sprinkled 400 times. In each sprinkling we take every spacelike pair with a near the top of the diamond ($\eta_a \in [-1.25, -1.05]$, $|x_a| \le 0.1$) and b at $x_b - x_a \in [0.1, 0.2]$, $|\eta_b - \eta_a| \le 0.05$ (9ℓ–17ℓ apart at $\rho = 10^4$). We keep each pair with a fixed probability, the same in every sprinkling, chosen so that P pairs remain on average (a position-based selection in the sense above), and pool the depth drift over them. We compare with $R = 2(f - \frac{1}{2})/(f_{\rm cont} - \frac{1}{2})$, obtained from the ordering fraction f of the whole diamond and calibrated on the exact continuum diamond ($f_{\rm cont} = 0.5173$). The errors are the

Table 3. Pooled curvature estimator $R_{\rm est} = -16\sum\beta/\sum\alpha$ over 1000 sprinklings ($H = 1$; window as in section 4: $N_C \geq 20$, down to the deepest c with $I(c,a)$ and $I(c,b)$ in the box). Pairs: inserted by hand at two target points; (nearest) the sprinkled elements nearest to them; (windows) all spacelike pairs of sprinkled elements in windows around them, about 20 per sprinkling (see text). Errors: bootstrap over sprinklings. $R_{\rm fit}^{\rm cont}$: the same fit on the exact continuum volumes of the same elements (de Sitter). Last column: standard deviation of single-pair values (de Sitter).

ρ	pairs	flat	de Sitter	$R_{\rm fit}^{\rm cont}$	single pair
4000	inserted	-0.03 ± 0.03	1.97 ± 0.11	2.07	3.5
	nearest	-0.07 ± 0.03	1.70 ± 0.11	2.07	3.6
	windows	0.00 ± 0.02	1.91 ± 0.10	2.08	3.5
10^4	inserted	-0.01 ± 0.02	1.99 ± 0.07	2.08	2.4
	nearest	-0.04 ± 0.02	1.84 ± 0.07	2.08	2.4
	windows	-0.01 ± 0.01	2.02 ± 0.06	2.08	2.4

spread between sprinklings.

As mean and standard deviation, the ordering fraction gives 2.05 ± 0.80 at $\rho = 4000$ and 1.99 ± 0.46 at $\rho = 10^4$. The pooled drift gives 1.67 ± 2.32 and 1.95 ± 1.34 for $P = 100$, against continuum values of 2.10 for the same fit: its spread is about three times larger. Pooling more pairs barely helps (standard deviation 2.78, 2.50, 2.32 and 2.27 for $P = 10$, 30, 100 and 300 at $\rho = 4000$), because pairs in one region share most of their common pasts. For the same reason, the ratio bias is not averaged away within one region: the estimate reads low by 0.43 ± 0.12 and 0.15 ± 0.07. It is the ratio that is biased, not the sums: forming the ratio from $\sum\beta$ and $\sum\alpha$ pooled over all 400 sprinklings gives 1.96 ± 0.12 and 2.04 ± 0.07. Drawing the pairs one at a time instead (a uniformly among the elements near the top, then b uniformly among those at the required offset) gives the same result within the errors in de Sitter space (1.63 and 1.92).

As a control we use a flat diamond of the same coordinate size, at a lower density, so that it holds the same number of elements as the de Sitter diamond (2.5×10^3 and 6.3×10^3). There the ordering fraction gives 0.01 ± 0.80 and 0.03 ± 0.50, and the pooled drift gives -0.16 ± 0.89 and -0.05 ± 0.54 (low by 0.16 ± 0.04 and 0.05 ± 0.03). The flat control is matched in element count, not in proper depth. In the same coordinates the deepest c lies about 1.85 below the pairs in flat space, against $\ln(3/1.15) \approx 0.96$ in de Sitter space, so the range of τ^2 over which the drift is fitted is about four times longer. That alone makes the flat drift less noisy, so the difference between the two geometries in the precision factors quoted next largely reflects this lever arm, not curvature. As a curvature estimator for a given region, the ordering fraction is therefore more precise, by a factor of about three in de Sitter space and slightly in flat space, and less biased. What the depth drift adds is geometric: the curvature appears as the failure of an exact flat-space identity, it is measured along a specific pair and depth rather than over a whole diamond, and equation (12) says where. Other counting routes use Ollivier–Ricci curvature [32], with which we have not compared.

5. A supporting diagnostic: the embedding test and its blind spot

Embedding. Given the squared intervals s_{ij} between M sampled elements, classical multidimensional scaling (MDS) with a Lorentzian Gram matrix reconstructs coordinates. Clough and Evans applied it to causal sets and proposed its eigenvalue spectrum as a test of dimension and manifoldlikeness [20, 21]. With $s_{ij} = (x_i - x_j)^\top \eta\,(x_i - x_j)$, double centering, $G = -\frac{1}{2} JSJ$ with $J = I - \mathbf{1}\mathbf{1}^\top/M$, gives $G = X\eta X^\top$ for the centered coordinates X. A $1+1$ Lorentzian set shows one large negative eigenvalue (time), one large positive eigenvalue (space), and others near zero; the two corresponding eigenvectors give coordinates up to a Lorentz transformation.

We supply MDS with counts only: $s = -2N(I)/\rho$ for related pairs, and equation (4) with the deepest c for spacelike pairs. On $M = 150$ elements of a flat sprinkling with $N = 2000$, positions are recovered to $0.43 \pm 0.07\,\ell$ across a region $\approx 38\ell$ wide. The best affine map from the reconstructed to the true coordinates, $x_{\text{true}} \approx A x_{\text{rec}} + b$, is a Lorentz transformation to within 3–10%, in the sense that the largest entry of $A^\top \eta A - \eta$ is 0.03–0.10 in absolute value. The two large eigenvalues are $\approx \mp 2.2$, and all others are below 0.1, in the coordinate units of the sprinkled region (a unit square in u, v).

Statistics and calibration. As test statistics we use intrinsic numbers only: r_+, the ratio of the second to the first positive eigenvalue; the share of the spectrum in one time and one space direction, $(|\lambda_-| + \lambda_+)/\sum_i |\lambda_i|$; and the *stress*, the rms misfit of the squared intervals by the best $1+1$ Minkowski embedding, relative to their rms. The sample consists of the $M = 40$ elements with the largest pasts, after removing elements with no common past with the rest. Figure 4(a) shows the calibration on flat universes (random two-dimensional orders). Stress and r_+ settle only for $N \gtrsim 100$ (stress 0.16, 0.10 and 0.08 at $N = 100$, 200 and 400), so the test is used only above this size. Curvature raises the stress: for de Sitter diamonds with fixed $N = 400$ it stays in the flat band up to $H\tau \approx 1$ and reaches 0.25 at $H\tau = 4$ [figure 4(c)]. The stress thus tests whether the order is a *flat* sheet, whereas section 4 measures its curvature.

Blind spot. Applied to Monte Carlo samples of the $\mathrm{e}^{-\beta S}$ model at $N = 400$ (section 6), this test passes every continuum-like state at that size (but not at $N = 800$; see section 6) and fails orders with 4 and 7 layers. It does not, however, reliably flag orders with about 10 layers: at $M = 40$ one of seven such orders lies in the flat band and two more lie just above it, and at $M = 80$ two others have stress 0.03, lower than flat orders. Two effects are responsible. First, the sample contains only the top one or two layers. Second, in a layered order most sampled pairs have zero squared interval: linked pairs have $N(I) = 0$, and spacelike pairs often have an empty strip, $N_A = 0$ or $N_B = 0$. MDS fits an all-zero matrix perfectly, and a bilayer produces exactly that matrix.

Null-pair fraction. We add the fraction z of sampled pairs with zero squared interval, taken as the larger of its values on the top sample and on a sample spread evenly in height (the longest chain below the element) over the upper 70% of the order. On sprinklings, z decreases with N (median 0.18, 0.12 and 0.09 at $N = 100$, 200 and 400) and does not respond to curvature: de Sitter diamonds at $H\tau = 4$ give 0.08–0.11.

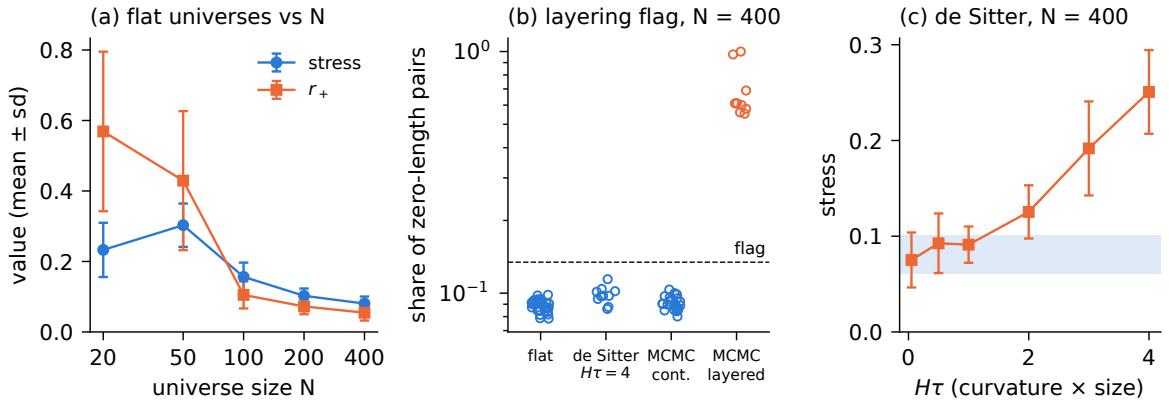


Figure 4. Calibration of the embedding test. (a) Flat universes (random two-dimensional orders, 10 per size, $M = 40$): stress and r_+ against the number of elements N. (b) Null-pair fraction z at $N = 400$ ($M = 80$) for flat orders, de Sitter diamonds at $H\tau = 4$, and final states of $e^{-\beta S}$ Monte Carlo chains (continuum-like, 30–39 layers; layered, 4–11 layers). Dashed: the flag threshold, 1.5 times the flat median. The 13 states with 2–3 layers are not shown. (c) Stress of de Sitter diamonds at fixed $N = 400$ (8 per height) against $H\tau$; shaded: the flat band (±1 standard deviation).

Layered orders give 0.4–1 [figure 4(b)]. We flag an order as layered when z exceeds 1.5 times the flat median at the same N. Orders with only 2–3 layers cannot be processed (too few elements share a common past), and the rule counts them as flagged.

What z adds. Layered orders are also detected by simpler statistics: the height (number of layers), the ordering fraction and the Myrheim–Meyer dimension [4, 5], interval abundances [7], and graph observables [41]. In the $e^{-\beta S}$ literature, the phases are characterized by the action, the ordering fraction, the height and interval abundances [27, 28, 16, 17, 18]. On our states the fixed-size ordering fraction separates the same groups as z (Appendix A). The null-pair fraction does not detect anything that these statistics miss. Its role is to make the embedding test usable on its own: without z the test passes some 10-layer orders, so a user who relied on the embedding spectrum alone would accept them. With z the test needs no second statistic, and it uses only quantities that the embedding already computes.

6. Applications

Sequential growth. In transitive percolation, the simplest classical sequential growth model [42], elements are born one at a time, each newborn element chooses each existing element as an ancestor with probability p, and the order is the transitive closure. As a sanity check of the diagnostics, we grew universes of $N = 100$ and 200 elements with $p = 0.05$, and with the non-covariant variant $p_n = 3/n$. Neither produces a $1 + 1$ spacetime. The ordering-fraction dimension drifts with scale, and the embedding test fails on every statistic: the share of the spectrum in one time and one space direction is ≈ 0.33, against 0.65 for flat orders at $N = 100$, and the stress is 0.64–0.78, against

Table 4. Embedding diagnostics on $e^{-\beta S}$ states at $N = 200$ ($\epsilon = 0.21$). f: ordering fraction; L: number of layers. The last column gives the range of the null-pair fraction z on fresh chains, for $\beta \leq 0.3$ (second row) and for $\beta \geq 0.6$ with 8–13 layers (third row).

$N = 200$	f	L	r_+	stress	z
random orders	0.50	24	0.07	0.10	0.10–0.15
$\beta = 0.05$	0.51	24	0.10	0.11	0.12–0.13
$\beta = 0.6$	0.59	9	0.39	0.28	0.39–0.48

0.10–0.12 at $N = 200$. It is not their null-pair fraction that flags them. Grown universes with $N = 1500$ also fail the propagator target (Appendix A). Covariant approaches to cosmology in growth models are discussed in [43, 44].

The two phases of the $e^{-\beta S}$ *model.* Weighting two-dimensional orders by $e^{-\beta S}$, with S the Benincasa–Dowker action, gives a first-order transition between a continuum-like and a layered phase [26, 27, 28]. We sampled the model with Metropolis swap moves at nonlocality parameter $\epsilon = 0.21$. Each update costs $O(N^2)$, since a swap changes two rows and columns of the causal matrix. Table 4 shows the diagnostics at $N = 200$. The states at $\beta \geq 0.6$ are metastable intermediates; the ground state at large β is a bilayer [45].

The states below are test configurations, representative of the two phases, not equilibrium samples (see the end of this section). We applied the flag to 44 final states of $N = 400$ chains run from random and from layered initial orders at $0.01 \leq \beta \leq 0.7$. All 22 continuum-like states (30–39 layers) pass. All 9 states with 4–11 layers are flagged ($z = 0.55$–1.0, against a threshold of 0.13). The 13 states with 2–3 layers cannot be processed and are counted as flagged. The fixed-n ordering fraction (Appendix A) separates the same groups; the propagator statistic misses one 4-layer state.

Because the flag was designed after these states had been seen, we tested it on held-out samples. On fresh $N = 200$ chains, 12 of 12 continuum-like states pass and 18 of 18 layered states ($\beta = 0.6$–1.5) are flagged. At $N = 800$ we used $M = 80$ (as for z at $N = 400$) and a threshold fixed in advance from flat orders: $z > 0.120$, 1.5 times the median $z = 0.080$ of 50 flat orders. None of 50 fresh flat orders and 20 de Sitter diamonds is flagged. We ran 28 chains ($\beta \in \{0.01, 0.05, 0.1, 0.2, 0.3, 0.45, 0.7\}$, 4 seeds each, random initial orders, 2000 sweeps), with the classes also fixed in advance: at least 20 layers, continuum-like; 4–19 layers, layered. This gives a small held-out sample of 8 layered and 20 continuum-like states. All 8 layered states (14–19 layers) are flagged ($z = 0.38$–0.57). Only one of them, with stress 0.072, lies in the flat band; this is the blind spot that z is meant to close, and the embedding spectrum alone would have passed it. Of the 20 continuum-like states (27–56 layers), 19 are not flagged ($z \leq 0.107$). The exception has 27 layers at $\beta = 0.3$ and is flagged with $z = 0.123$, just above the threshold. Its acceptance rate (5%, against 53–54% for the other three chains at $\beta = 0.3$, which reached 43–50 layers) suggests that it is partly layered, but by the classes fixed in

advance it is a false positive. For context only: random two-dimensional orders of 800 elements have 52 layers on average (standard deviation 2.5), so this state has about half the flat height, whereas the continuum-like states at $N = 400$ had 30–39 layers against a flat mean of 36. Three of the 19 unflagged states have a stress slightly above the flat band at $N = 800$ (0.086–0.110, against the flat mean plus two standard deviations, 0.081), as does the flagged one (0.134). At this size, therefore, the stress alone would also reject some continuum-like states.

The $N = 400$ states are test inputs, not equilibrium samples: chains started from layered orders accept $\leq 0.1\%$ of the moves. They bracket the transition only loosely: chains from layered starts become continuum-like at $\beta \leq 0.05$, and chains from random starts become layered at $\beta \geq 0.45$. The bracket $0.05 < \beta_c(400) < 0.45$ is compatible with the finite-size scaling $\beta_c \approx 1.66/(\epsilon^2 N) = 0.094$ of [28], but it is too loose to test it.

7. Discussion

What is new. The overlap regions are those of [10], the embedding test is that of [20], the null-pair fraction is related to [7], and the Laguerre identity is classical combinatorics. To our knowledge, the following are new: the rectangle identity for general c, equation (3), with the bias-corrected estimator and its comparison with existing constructions, including a counting version of the Hu–Meyer–Meyer construction; the depth drift, its closed form in de Sitter space, equation (11), and its leading form in FRW spacetimes, equation (12), with the curvature estimator built on it; the application of the Laguerre identity to Johnston's propagator, equation (A.2); and the calibration of the embedding test against size, curvature and layering. Our literature search (October 2026) covered the papers citing [10, 20, 24, 25]. The rectangle identity is elementary enough that it may appear in passing in work we did not reach.

What the curvature result is. Since the ordering fraction of a region is the more precise estimate of its curvature, the curvature result of this paper is analytic and conceptual: an exact flat-space identity whose depth-dependent failure has a closed form, equation (11), and a leading-order meaning, equation (12). Whether the locality of the depth drift, a curvature along one pair and one depth, makes it the better estimator where curvature varies across a region is untested. A methodological point is of more general use. Estimators evaluated on sprinkled elements chosen as the nearest to given points, or one per window, inherit a Palm-type selection bias, because the choice depends on the neighborhood that the estimator counts. Summing over all elements or pairs of a position-based selection, as the Mecke formula requires, removes it.

Limitations. The exact results are specific to $1 + 1$ dimensions. The rectangle identity has no direct analogue in higher dimensions, where the overlap of two intervals is not an interval. The distance estimator is the most accurate only below about 10ℓ; at larger separations the counting version of the Hu–Meyer–Meyer construction is better. The curvature estimator must be pooled over all pairs of a position-based selection and over many independent regions; within one region the ratio stays biased (see below).

Single pairs are dominated by noise and their ratio is biased, and selecting the elements nearest to given points, or one element per window, biases it low. The target of the estimator is the same fit on the continuum volumes, not R itself. Because the window reaches $H\tau \approx 1$, the τ^4 term of equation (11) makes this target 2.07–2.08 instead of 2 in our de Sitter box, and 2.03–2.04 in the convex region. This is a systematic error of the method, of +2 to +4%, which depends on the size of the region and is not known in practice; capping the window lowers it at the cost of noise (section 4). In FRW spacetime the leading-order prediction of equation (12) falls short of this target by 26–35% over realistic windows. Within one region the pairs share elements, so pooling more than a few tens of them gains little, and the estimate retains a low ratio bias of 0.05–0.43 at the densities used. For the curvature of a given region the ordering fraction is more precise, by a factor of about three in de Sitter space, and less biased. The estimator was tested on de Sitter space and on one FRW spacetime. The embedding thresholds are empirical and must be recalibrated at each N. On a small held-out sample at $N = 800$ (8 layered and 20 continuum-like states), the layered flag caught every layered state but also one of the 20 continuum-like states, just above its threshold. The stress put that state (0.134) and 3 of the 19 unflagged ones above the flat band, so at this size the full embedding test (spectrum, stress and z) rejects 4 of the 20 continuum-like states. The $\mathrm{e}^{-\beta S}$ states used are test inputs, not equilibrium samples. The Laguerre target is exact only for two-dimensional orders, and as a manifoldlikeness diagnostic it is not stronger than the ordering fraction.

Outlook. Since the Hu–Meyer–Meyer construction is the more accurate one at large separations, a depth drift built on it may be less noisy there; a combined distance estimator would switch between the two near 10ℓ. A fixed-n propagator target in $3+1$ dimensions would not be exact. Johnston's $3+1$ particle hops along links only [24], so a path requires every skipped sub-interval to be empty, and $E[K \mid n]$ becomes a sum of $(k-1)$-point integrals with factors $(1 - W/V)^{n-k+1}$. Its first term can be evaluated at large n: the expected number of two-hop link paths tends to 3π. This is what Johnston's amplitudes require to reproduce the short-distance limit $m^2/8\pi$ of the continuum mass term [24], so it is a consistency check rather than a new result; a Monte Carlo evaluation approaches it as $3\pi(1-1.6/\sqrt{n})$ at fixed n. A $3+1$ target would therefore be a computed table rather than a formula.

Data availability statement

The code that produces every number and figure in this paper, and the research notebook in which the work was carried out, are openly available at `https://github.com/presindev/causet-counting`, and the version used for this paper is archived at doi:[Zenodo DOI]. Figures are produced by `sims/fig_part1.py`; the baseline comparison by `sims/r31_rect_baseline.py` and, with the Hu–Meyer–Meyer construction and the ablations, by `sims/r36_rect_ablation_hmm.py`; the curvature estimator (window, pooled form, sprinkled pairs, Benincasa–Dowker and

chain-abundance comparison, bin correlations) by `sims/r37_curvature_referee.py`; the sprinkled-pair selections, the causally convex region, the FRW test with the pooled estimator and the first comparison with the ordering fraction in one region by `sims/r40_curvature_round2.py`; the one-region comparison with a position-based selection and the flat control at matched element count by `sims/r41_curvature_round3.py`; equation (12) by `sims/r38_drift_kernel.py`; the null-pair fraction by `sims/r8_detector_spread.py`; the curvature fit and the FRW check by `sims/r9_drift_fit_frw.py`; the $e^{-\beta S}$ states by `sims/n400_crystal.py`; the propagator statistics by `sims/b_johnston_n400.py` and `sims/r32_laguerre_vs_fraction.py`; and the $N = 800$ test by `sims/r27_n800_detector.py`, with the heights of random orders quoted with it by `sims/r42_flat_height.py`.

Acknowledgments

This work was funded by SapereIA. Simulations were run on the Nyx cluster of the Institute of Cosmos Sciences of the University of Barcelona (ICCUB). An AI system (Claude Opus 5.5, Anthropic) was used for literature search, derivations, code and drafting, under the author's direction and checks; the procedure is described in [47]. The author takes full responsibility for the content.

Appendix A. A density-free target for Johnston's propagator

This appendix records a combinatorial identity for Johnston's propagator. As a diagnostic of manifoldlikeness it is not stronger than the ordering fraction (see the end of this appendix); its use is as an exact, density-free check of implementations of the chain sum.

The propagator. In Johnston's model a particle hops along the order with amplitude a per hop and b per intermediate stop [24]. With C the causal matrix ($C_{xy} = 1$ if $x \prec y$) and $\Phi = aC$,

$$K = \Phi + b\,\Phi^2 + b^2\Phi^3 + \cdots = \Phi\,(I - b\Phi)^{-1}, \tag{A.1}$$

with $a = \frac{1}{2}$ and $b = -m^2/\rho$ in $1+1$ dimensions. $K(x,y)$ is a weighted count of chains, $K(x,y) = \frac{1}{2}\sum_k c_k(x,y)\,(-\mu/2)^k$, where c_k is the number of k-element chains strictly inside $I(x,y)$ and $\mu = m^2/\rho$. In a sprinkled interval of volume $V = \tau^2/2$ the expected number of k-chains is $(\rho V)^k/(k!)^2$, so that $\langle K\rangle = \frac{1}{2}\sum_k[-(m\tau/2)^2]^k/(k!)^2 = \frac{1}{2}J_0(m\tau)$, the retarded Klein–Gordon propagator [24].

Fixed interval size. K depends only on the order and on μ. In a uniformly random two-dimensional order, the n elements inside an interval again form a uniformly random two-dimensional order, and any k of them form a chain with probability $1/k!$ (the two

linear orders must agree on them). Hence $E[c_k \,|\, n] = \binom{n}{k}/k!$ [34], and

$$\begin{aligned} E[K \,|\, n] &= \frac{1}{2}\sum_{k=0}^{n}\binom{n}{k}\frac{(-\mu/2)^k}{k!} = \frac{1}{2}L_n\Big(\frac{\mu}{2}\Big) \\ &\approx \frac{1}{2}\mathrm{e}^{\mu/4}J_0\Big(\sqrt{2\mu(n+\tfrac{1}{2})}\Big), \end{aligned} \tag{A.2}$$

a Laguerre polynomial in the interval size; the Bessel form is the Hilb-type asymptotic for large n at fixed μ [35]. Averaging over Poisson-distributed n with mean $\lambda = \rho\tau^2/2$, using $\sum_n \mathrm{e}^{-\lambda}\frac{\lambda^n}{n!}\binom{n}{k} = \frac{\lambda^k}{k!}$, recovers $\frac{1}{2}\sum_k(-\mu\lambda/2)^k/(k!)^2 = \frac{1}{2}J_0(m\tau)$: Johnston's result is the Poisson average of equation (A.2). The fixed-n form requires neither the density nor coordinates: every related pair of a causal set supplies its n and its K, and equation (A.2) is the target. The identity itself is combinatorial: at $\mu = -2$, and summed over permutations, it counts all increasing subsequences, $n!\,L_n(-1)$ [36, 37]. A fixed-n test also exists for maximal chains, whose length distribution in a two-dimensional interval of n elements measures dimension and embeddability [38]; equation (A.2) instead sums all chains with the propagator's weights. Path-sum studies of the propagator average over sprinklings [39, 40].

Tests. Figure A1 bins all related pairs of an $N = 400$ order by n and compares $\langle K\rangle$ with equation (A.2) at $\mu = 0.05$. As a statistic we use the rms deviation S_K of K from the target over all related pairs with $n \geq 2$, and compare it with its distribution over 40 random orders ($S_K = 0.0158 \pm 0.0018$). We apply it to 60 final states of $N = 400$ chains of the $\mathrm{e}^{-\beta S}$ model (section 6; the 44 states used there and 16 from chains completed later): 28 continuum-like states (30–39 layers), 14 states with 4–11 layers, and 18 with 2–3 layers, which contain too few large intervals for the statistic. Continuum-like states lie within 3 standard deviations of the random orders. States with 9–11 layers deviate by 39–62 standard deviations. In 4- to 7-layer crystals no interval contains more than a few elements, so only small n, where the target is close to $\frac{1}{2}$, is probed. Their deviations range from -9 to $+25$ standard deviations, and one 4-layer state lies within 1.2 standard deviations of the random orders. S_K alone therefore does not flag every crystal. Sequentially grown universes ($N = 1500$) deviate strongly, one of them ($p = 0.02$) reaching $|K| \approx 10$ for individual pairs, whereas the target is bounded by $\frac{1}{2}$. Since such orders are not two-dimensional orders, this only confirms that the target is not met by arbitrary orders.

Relation to the ordering fraction. Expanding equation (A.2) in μ, the first nontrivial term is c_2, the number of related pairs inside the interval, whose fixed-n mean $\binom{n}{2}/2$ is the ordering fraction behind the Myrheim–Meyer dimension [4, 5]. We therefore compared S_K with the analogous statistic S_f, the rms deviation of $c_2/\binom{n}{2}$ from $\frac{1}{2}$, on the same states. The layered states deviate from random orders by 50–58 standard deviations in S_f, comparable to S_K, and all six 4- to 7-layer crystals deviate by more than 100, including the one that S_K misses. As a diagnostic of manifoldlikeness, the propagator target is therefore not stronger than its first term. Its use lies elsewhere: it is an exact, density-free statement about the propagator itself, against which an implementation

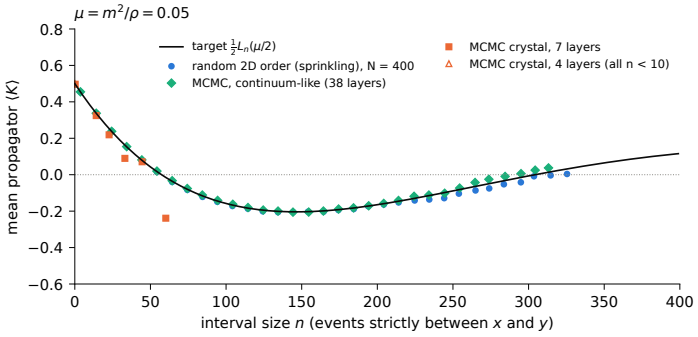


Figure A1. Johnston's propagator averaged over related pairs with interval size n, against $\frac{1}{2}L_n(\mu/2)$, at $\mu = 0.05$ and $N = 400$. At large n few pairs remain and the bins are noisy. In the 4-layer crystal every interval has fewer than 10 elements, so it contributes a single point near $n = 0$.

of the chain sum can be checked pair by pair on any two-dimensional order, without knowledge of the density.

Appendix B. Other approaches tested

- *Fixed-n propagator as a curvature estimator.* In curved $1 + 1$ spacetime a term $(R\tau^2/48)J_2(m\tau)$ appears in the propagator [46]. In a de Sitter diamond conditioned on n elements, the order depends only on $(n, H\tau)$, so a fixed-n curvature test exists. We find that the fraction of related pairs rises by $\approx 0.027\,(H\tau)^2$ for $n = 10$–60, and that the fixed-n propagator is no more sensitive than that fraction alone: a 3σ detection at $n = 30$, $H\tau = 1$ requires 175 intervals in either case. The fixed-n propagator therefore detects curvature only through the known chain-abundance route [6].
- *Embedding test on small orders.* We first applied the test to exhaustive ensembles of small orders ($N \leq 12$). At these sizes its statistics do not separate flat from layered orders; they settle only for $N \gtrsim 100$ [figure 4(a)], so we moved to Monte Carlo samples at $N \geq 200$.
- *Hysteresis in the* $\mathrm{e}^{-\beta S}$ *chains.* At $N = 400$, chains started from random and from layered orders at the same β end in different phases. We first read this as hysteresis of the transition. However, the chains from layered starts accept $\leq 0.1\%$ of the moves and are effectively frozen, so the difference reflects the sampler rather than the model.
- *Two-dimensional estimator on* $2 + 1$ *sprinklings.* Fed into the embedding, equation (4) applied to a $2+1$ sprinkling (with the $2+1$ density) produces two large positive eigenvalues. The formula is not valid there, so this is at most a qualitative indication that the spectrum counts spatial dimensions.

References

[1] Bombelli L, Lee J, Meyer D and Sorkin R D 1987 Space-time as a causal set *Phys. Rev. Lett.* **59** 521

[2] Surya S 2019 The causal set approach to quantum gravity *Living Rev. Relativ.* **22** 5 (arXiv:1903.11544)

[3] Brightwell G and Gregory R 1991 Structure of random discrete spacetime *Phys. Rev. Lett.* **66** 260

[4] Myrheim J 1978 Statistical geometry *CERN Report* TH-2538

[5] Meyer D A 1988 The dimension of causal sets *PhD Thesis* Massachusetts Institute of Technology

[6] Roy M, Sinha D and Surya S 2013 Discrete geometry of a small causal diamond *Phys. Rev.* D **87** 044046 (arXiv:1212.0631)

[7] Glaser L and Surya S 2013 Towards a definition of locality in a manifoldlike causal set *Phys. Rev.* D **88** 124026 (arXiv:1309.3403)

[8] Benincasa D M T and Dowker F 2010 The scalar curvature of a causal set *Phys. Rev. Lett.* **104** 181301 (arXiv:1001.2725)

[9] Dowker F and Glaser L 2013 Causal set d'Alembertians for various dimensions *Class. Quantum Grav.* **30** 195016 (arXiv:1305.2588)

[10] Boguñá M and Krioukov D 2024 Measuring spatial distances in causal sets via causal overlaps *Phys. Rev.* D **110** 024008 (arXiv:2401.17376)

[11] Rideout D and Wallden P 2009 Spacelike distance from discrete causal order *Class. Quantum Grav.* **26** 155013 (arXiv:0810.1768)

[12] Major S, Rideout D and Surya S 2006 Spatial hypersurfaces in causal set cosmology *Class. Quantum Grav.* **23** 4743 (arXiv:gr-qc/0506133)

[13] Major S, Rideout D and Surya S 2007 On recovering continuum topology from a causal set *J. Math. Phys.* **48** 032501 (arXiv:gr-qc/0604124)

[14] Eichhorn A, Surya S and Versteegen F 2019 Induced spatial geometry from causal structure *Class. Quantum Grav.* **36** 105005 (arXiv:1809.06192)

[15] Khetrapal S and Surya S 2013 Boundary term contribution to the volume of a small causal diamond *Class. Quantum Grav.* **30** 065005 (arXiv:1212.0629)

[16] Glaser L and Surya S 2016 The Hartle–Hawking wave function in 2D causal set quantum gravity *Class. Quantum Grav.* **33** 065003 (arXiv:1410.8775)

[17] Glaser L 2018 The Ising model coupled to 2d orders *Class. Quantum Grav.* **35** 084001 (arXiv:1802.02519)

[18] Glaser L 2023 Computer simulations of causal sets arXiv:2306.09904

[19] Hu C A, Meyer D A and Meyer E J Q 2025 Reconstructing Minkowski geometry from causal separations *J. Math. Phys.* **66** 122502 (arXiv:2601.03280)

[20] Clough J R and Evans T S 2017 Embedding graphs in Lorentzian spacetime *PLoS One* **12** e0187301 (arXiv:1602.03103)

[21] Clough J R 2017 Causal structure in networks *PhD Thesis* Imperial College London (hdl:10044/1/50162)

[22] Henson J 2006 Constructing an interval of Minkowski space from a causal set *Class. Quantum Grav.* **23** L29 (arXiv:gr-qc/0601069)

[23] Johnston S 2025 Simpler embeddings of causal sets into Minkowski spacetime *Phys. Rev.* D **111** 106020 (arXiv:2502.09701)

[24] Johnston S 2008 Particle propagators on discrete spacetime *Class. Quantum Grav.* **25** 202001 (arXiv:0806.3083)

[25] Johnston S 2009 Feynman propagator for a free scalar field on a causal set *Phys. Rev. Lett.* **103** 180401 (arXiv:0909.0944)

[26] Brightwell G, Henson J and Surya S 2008 A 2D model of causal set quantum gravity: the emergence of the continuum *Class. Quantum Grav.* **25** 105025 (arXiv:0706.0375)

[27] Surya S 2012 Evidence for the continuum in 2D causal set quantum gravity *Class. Quantum Grav.*

29 132001 (arXiv:1110.6244)

[28] Glaser L, O'Connor D and Surya S 2018 Finite size scaling in 2d causal set quantum gravity *Class. Quantum Grav.* **35** 045006 (arXiv:1706.06432)

[29] Bada R 2026 Spatial embedding of $3+1$ causal sets: low-count spacelike pairs and a likelihood fit to interval counts (companion paper, to be submitted)

[30] Baik J, Deift P and Johansson K 1999 On the distribution of the length of the longest increasing subsequence of random permutations *J. Am. Math. Soc.* **12** 1119

[31] Cunningham W J 2018 Inference of boundaries in causal sets *Class. Quantum Grav.* **35** 094002 (arXiv:1710.09705)

[32] Barton J, Borza S and Röhrig J 2026 Ollivier–Ricci curvature for causal sets arXiv:2606.04910

[33] Last G and Penrose M 2017 *Lectures on the Poisson Process* (Cambridge: Cambridge University Press) ch 4

[34] Pinsky R G 2006 Law of large numbers for increasing subsequences of random permutations *Random Struct. Algorithms* **29** 277 (arXiv:math/0407353) equation (1.1)

[35] Szegő G 1939 *Orthogonal Polynomials* (American Mathematical Society Colloquium Publications vol 23) (New York: American Mathematical Society)

[36] Lifschitz V and Pittel B 1981 The number of increasing subsequences of the random permutation *J. Comb. Theory* A **31** 1

[37] OEIS Foundation The On-Line Encyclopedia of Integer Sequences, sequence A002720 https://oeis.org/A002720

[38] Aghili M E, Bombelli L and Pilgrim B B 2018 Path length distribution in two-dimensional causal sets *Eur. Phys. J.* C **78** 744 (arXiv:1805.07312)

[39] Shuman S 2024 Path sums for propagators in causal sets *Phys. Rev.* D **109** 046008 (arXiv:2307.08864)

[40] Kastrati A and Hinrichsen H 2026 Numerical evaluation of the causal set propagator in 2D anti-de Sitter spacetime *Phys. Rev.* D **113** 065003 (arXiv:2504.12919)

[41] Eichhorn A, Mack H, Le K T and Wagner F 2026 Charting causal set configuration space with graph observables arXiv:2605.27514

[42] Rideout D P and Sorkin R D 2000 Classical sequential growth dynamics for causal sets *Phys. Rev.* D **61** 024002

[43] Martín X, O'Connor D, Rideout D P and Sorkin R D 2001 On the renormalization transformations induced by cycles of expansion and contraction in causal set cosmology *Phys. Rev.* D **63** 084026

[44] Ahmed M and Rideout D 2010 Indications of de Sitter spacetime from classical sequential growth dynamics of causal sets *Phys. Rev.* D **81** 083528

[45] Cunningham W J and Surya S 2020 Dimensionally restricted causal set quantum gravity: examples in two and three dimensions *Class. Quantum Grav.* **37** 054002 (arXiv:1908.11647)

[46] Nomaan X, Dowker F and Surya S 2017 Scalar field Green functions on causal sets *Class. Quantum Grav.* **34** 124002 (arXiv:1701.07212)

[47] Bada R 2026 Two causal-set papers written with an AI research partner: methodology note, released with the code at https://github.com/presindev/causet-counting